

QRCS Corporation

# Hierachal Key Distribution System Analysis

**Title:** Implementation Analysis of the Hierachal Key Distribution System (HKDS)

**Author:** John G. Underhill

**Institution:** Quantum Resistant Cryptographic Solutions Corporation (QRCS)

**Date:** November 2025

**Document Type:** Technical Cryptanalysis Report

**Revision:** 1.0

## Chapter 1: Introduction and Problem Statement

This paper analyzes the Hierachal symmetric Key Distribution System (HKDS), a two-key, server-assisted symmetric key distribution protocol intended to replace DUKPT in high-scale payment environments. HKDS couples a per-device Embedded Device Key (EDK) with server-derived ephemeral token keys and uses Keccak-family primitives (SHAKE as PRF/XOF and KMAC for authentication). The specification positions HKDS as a scalable, authenticated alternative supporting 128/256/512-bit security levels and designed for point-of-sale, ATM, and similar deployments in financial services.

### 1.1 Context and motivation

Payment ecosystems now depend on vast fleets of remote terminals communicating with central processors. The HKDS specification identifies distributed key management as the core scaling bottleneck and frames DUKPT as a legacy architecture whose cost and scalability no longer match present traffic and device counts. HKDS is proposed as a replacement that preserves low device cost while enabling authenticated operation and higher security margins.

### 1.2 Informal protocol overview

Each deployment maintains a Master Derivation Key (MDK) on the server, comprising a Base Derivation Key (BDK), a Secret Token Key (STK), and a 32-bit key identifier (KID). The client device holds a unique EDK derived from the BDK and a Key Serial Number (KSN) containing KID, a device identity (DID), and a transaction counter (KTC).

On initialization and whenever the transaction-key cache is exhausted, the client transmits a token request containing its KSN. The server authenticates the request context, derives a per-request secret token from STK and a customization string (CTOK), and returns an encrypted token (ETOK) with a 16-byte KMAC tag. The client verifies the tag using a token-MAC string (TMS/TSM) derived from KSN and the formal MAC name, decrypts ETOK with an EDK-bound token-encryption key, and regenerates its transaction-key cache by seeding SHAKE with  $\text{TOK} \parallel \text{EDK}$ .

During message protection, the client selects two cache entries: one to XOR-encrypt the 16-byte PIN (stream-cipher-style), and a second to key KMAC over the ciphertext (encrypt-then-MAC). The server reconstructs the same cache from (EDK, token) and verifies/decrypts using the corresponding entries.

HKDS uses SHAKE and KMAC with fixed parameters chosen at implementation time. Protocol headers identify message type, sequence, and length, but do not negotiate or signal alternate PRF modes.

### 1.3 Design goals

The specification asserts three primary goals: (1) forward security of used transaction keys, (2) predictive resistance under client state compromise, and (3) bounded server recomputation to recover transaction keys. It further emphasizes a two-key construction (EDK plus server-derived token) to enable periodic entropy refresh without updating embedded keys, and to permit unbounded transaction-key generation subject to periodic token refresh.

### 1.4 Cryptographic primitives and implementation surface

HKDS uses SHAKE as the PRF/XOF for EDK derivation, token derivation, and cache generation, and uses KMAC for token authentication and message authentication in the encrypt-then-MAC channel. The accompanying Keccak implementation defines the SHAKE/KMAC domains, rates, state size, and round controls needed by the construction. We will analyze security claims under the assumption that SHAKE and KMAC instantiated per SP 800-185/FIPS 202 behave as standard Keccak-based primitives.

### 1.5 Threat surface at a glance

From the specification's operational flows, the principal adversarial vantage points include: passive observation of messages and KSNs; active manipulation of token

exchanges or cache synchronization; client compromise revealing current cache entries or partial state; and server compromise affecting MDK secrecy. HKDS aims to bind token and cache derivations to DID, CTOK, and counters to ensure per-device, per-request uniqueness and to allow bounded server recomputation. Formal adversary models and reduction targets will be stated in Chapter 3.

## 1.6 Scope of this analysis

This work evaluates correctness of the derivation flows, authenticity guarantees in the token and message channels, key-separation and misuse resistance within the two-key design, cache synchronization robustness, and the impact of SHAKE/KMAC parameterization on security margins and performance. We will restrict claims to what is supported by the specification and provided source, and will explicitly label conjectures. Implementation-specific behavior (e.g., SHAKE rates, domain identifiers, round parameters) will be checked against the included Keccak code.

# Chapter 2: Model, Assumptions, and Security Objectives

## 2.1 System model

The HKDS architecture comprises two communicating entities: the **Client Device** (terminal) and the **Key Server** (issuer). Each device  $i$  has a unique 32-bit Device Identifier ( $DID_i$ ) and an Embedded Device Key ( $EDK_i$ ), derived during initialization from the Base Derivation Key (BDK) and  $DID_i$  through SHAKE or KMAC. The Key Server maintains the Master Derivation Key (MDK) structure, composed of BDK, a Secret Token Key (STK), and a Key Identifier (KID).

At runtime, the Client produces a Key Serial Number (KSN) defined as  $KSN = (KID \parallel DID_i \parallel KTC)$ , where KTC is a monotonic transaction counter incremented per key use. The KSN authenticates the derivation state and transaction index, forming the primary synchronization element between Client and Server.

The Server derives a per-request token T from STK and a customization string  $CTOK$  that encodes  $DID_i$ , KID, and other contextual data. The token is transmitted as an encrypted token ETOK under a key derived from  $EDK_i$  and authenticated via KMAC using a Token Signature Message (TSM).

Both Client and Server maintain a transaction-key cache Kcache containing a sequence of subkeys ( $K_1, K_2, \dots, K_n$ ) generated by applying SHAKE to the concatenation of EDK<sub>i</sub> and token T. These subkeys are consumed sequentially to protect each transaction.

In an operational deployment, all session authentication and integrity checks are handled via KMAC tags computed under derived cache keys. The Client uses one cache key for XOR-encryption of the PIN or message body, and a second key for KMAC authentication. The Server verifies the tag and decrypts by regenerating the same cache from its stored MDK, STK, and the KSN supplied by the Client.

## 2.2 Security assumptions

The following assumptions are drawn from the specification and the Keccak-based implementation provided:

### 1. Cryptographic primitives.

SHAKE and KMAC are modeled as ideal extendable-output and keyed hash functions respectively, with collision, preimage, and distinguishing advantages negligible in  $2^{-b}$  for security strength  $b = 128, 256$ , or  $512$ . The internal Keccak permutation is assumed to follow the FIPS 202 round structure without omission or bias.

### 2. Randomness and counters.

Each EDK is unique and uniformly distributed across its key-space. The KTC counter is monotonic and never wraps within the operational lifetime of the device. Server tokens are independently derived for each request, and token identifiers (CTOK values) are non-repeating per DID.

### 3. Adversary limitations.

The adversary has full visibility of public values (KSN, ciphertexts, MACs) and may inject, replay, or modify messages but cannot compromise both EDK<sub>i</sub> and STK simultaneously. The Server is considered honest-but-curious unless stated otherwise. Side-channel leakage and hardware tampering are excluded from the primary model.

### 4. Implementation correctness.

The Keccak reference implementation is presumed faithful to specification, using standard padding ("pad10\*1") and domain separation constants. No timing-dependent or key-dependent control flow exists in the critical path.

## 2.3 Security objectives

HKDS declares three explicit objectives:

### Forward secrecy.

Once a transaction key  $K_t$  has been used and the corresponding token and cache material have been erased, compromise of the remaining client state must not allow recovery of  $K_t$ . This requires the SHAKE stream used to generate the cache to be irreversible with respect to its internal position and to depend jointly on the token and the embedded device key.

### Predictive resistance.

Given  $EDK_i$  and previously derived keys, the adversary must not predict future cache entries without knowledge of token  $T$ , since token entropy derives from STK unknown to the Client.

### Server recoverability.

The Server must be able to regenerate any Client's transaction key on demand, using MDK and KSN as inputs. This imposes determinism across the key-derivation pipeline, ensuring that identical ( $EDK_i, T, KTC$ ) inputs yield identical cache outputs.

## 2.4 Derived requirements

From the objectives above, the following operational constraints arise:

- Key-derivation functions must be domain-separated to avoid cross-protocol collisions between EDK, token, and cache derivations.
- The encryption and authentication keys derived from cache material must not overlap in bits; internal expansion via SHAKE suffices if the bitstream offsets are non-overlapping.
- Token replay or duplication must be detectable by inclusion of KTC and DID in the CTOK or TSM components.
- The Server must reject stale or inconsistent KSN values; Client caches must be invalidated upon resynchronization failures.

## 2.5 Model summary

The analytical model thus comprises a two-party, stateful symmetric key distribution system under a standard indistinguishability-under-chosen-message (IND-CPA) and unforgeability (UF-CMA) framework, instantiated with Keccak-based primitives. The remainder of this paper evaluates whether the HKDS design and its provided implementation meet these objectives under the stated assumptions and whether any operational or structural weaknesses undermine these goals.

## Chapter 3: Formal Specification and Derivation Functions

### Formal Definition of the HKDS Construction

This chapter defines the HKDS mechanism precisely as implemented and specified. All functions, inputs, and concatenation orders follow the HKDS specification and the reference implementation without modification. No idealized domain tags or auxiliary labels are introduced. All cryptographic operations rely on Keccak-based primitives used in their standard SP 800-185 form.

Let  $\text{SHAKE}_p$  denote the extendable output function with rate  $p$ . Let  $\text{KMAC}_p$  denote the message authentication function defined in SP 800-185 using Keccak with fixed output length with a rate of  $p$ . All concatenations are literal byte concatenations.

HKDS consists of four deterministic processes. These are device key derivation, token generation, encrypted token transport with integrity protection, and transaction key cache generation.

The shared server state contains two private roots. The base derivation key BDK is used to provide a persistent device specific secret. The server token key STK is used to derive short lived tokens that drive forward secrecy and predictive resistance. Each client instance stores its device identifier DID, its key serial number KSN, the derived device key EDK, and a transaction key cache TKC.

The construction is defined by the following components.

#### 3.1 Embedded Device Key Derivation

Each device is provisioned with a device identifier DID. The server maintains a private root key BDK. The embedded device key EDK is deterministically derived as

$EDK = \text{SHAKE}_\rho(DID \parallel BDK)$ .

This operation binds the device identity to the server root. The derivation is injective with respect to DID and BDK and produces a uniformly distributed secret key of the required length.

The client stores EDK permanently after provisioning. The server recomputes EDK whenever required using the same inputs.

### 3.2 Customization Block for Token Derivation

Each token is bound to a specific device and a specific region within the transaction key cache. This is determined by CTOK. The customization block is the concatenation

$CTOK = TRC \parallel FN \parallel DID$ .

Here TRC is the transaction region code, computed as the integer division of the device transaction counter KTC by the cache size. FN is the formal name string of the HKDS instance as specified in the implementation. DID is the device identifier.

CTOK uniquely identifies the cache region and device identity for the token.

### 3.3 Token Derivation

The server maintains a private server token key STK. For a given CTOK, the server constructs the token as

$TOK = \text{SHAKE}_\rho(CTOK \parallel STK)$ .

TOK is a secret value that the client uses to generate the transaction key cache. The server uses only deterministic recomputation and never persists TOK.

### 3.4 Encrypted Token Transport

The server transmits TOK to the client in encrypted form to prevent token disclosure.

The encryption pad is generated as

$Kpad = \text{SHAKE}_\rho(CTOK \parallel EDK)$ .

The encrypted token is then

$ETOK = TOK \oplus Kpad$ .

This construction provides confidentiality for the token under the assumption that SHAKEp behaves as a pseudorandom function and that EDK remains unknown to the adversary.

### 3.5 Token Message Authentication Tag

Each encrypted token is authenticated to prevent tampering or substitution. The authentication uses KMACp with the embedded device key EDK as the key and a customization string TMS that binds the authentication to the device identity and the MAC name. The customization string is

$$TMS = KSN \parallel MACname.$$

The token authentication tag is

$$TAGt = KMACp(EDK, ETOK, custom = TMS).$$

The client verifies TAGt using its stored EDK, its known KSN, and the MAC name. Successful verification ensures authenticity and freshness of the token with respect to the server.

### 3.6 Token Recovery

The client recovers the token after verifying TAGt. It deterministically recomputes the encryption pad

$$Kpad = SHAKEp(CTOK \parallel EDK)$$

and decrypts

$$TOK = ETOK \oplus Kpad.$$

This operation produces the same TOK value held by the server for this transaction region.

### 3.7 Transaction Key Cache Derivation

Both client and server generate the transaction key cache deterministically from TOK and EDK using

$$TKC = SHAKEp(TOK \parallel EDK).$$

The output is parsed into a sequence of fixed size blocks

$TKC = (K_1, K_2, \dots, K_n)$

where  $n$  equals the HKDS cache size and each block  $K_i$  has length equal to `HKDS_MESSAGE_SIZE`.

The  $i$ -th transaction key used for message processing corresponds to an index determined by the evolving transaction counter. Only sequential use of keys is permitted.

### 3.8 Message Encryption

Let  $K_i$  denote the first unused cache entry. Let  $M$  denote the plaintext of `HKDS_MESSAGE_SIZE` bytes. The ciphertext is produced as

$$C = M \oplus K_i.$$

This operation is equivalent to a one-time pad under the assumption that  $K_i$  is indistinguishable from random and is never reused.

### 3.9 Message Authentication

Let  $K_{i+1}$  denote the next unused cache entry. Let  $A$  denote optional associated data. The message authentication tag is

$$TAG = KMACp(K_{i+1}, C \parallel A, \text{custom} = \text{empty}).$$

The server verifies  $TAG$  by recomputing  $K_i$  and  $K_{i+1}$  using the reconstructed token and transaction counter.

### 3.10 Server Side Deterministic Reconstruction

For every client message, the server recomputes all required values without storage of long term per device state beyond BDK, STK, and KTC. Given DID, KSN, and the current counter, the server deterministically performs

1.  $EDK = \text{SHAKE}\rho(DID \parallel BDK)$
2.  $CTOK = \text{TRC} \parallel \text{FN} \parallel \text{DID}$
3.  $TOK = \text{SHAKE}\rho(CTOK \parallel STK)$
4.  $TKC = \text{SHAKE}\rho(TOK \parallel EDK)$

and then selects the appropriate  $K_i$  and  $K_{i+1}$  based on the counter modulo the cache size.

This ensures decryptability, verifiability, and forward secrecy.

## Chapter 4: Mathematical Formulation and Security Properties

### Security Assumptions and Domain Separation

This chapter defines the security assumptions underlying HKDS and formalizes the manner in which key separation and domain isolation arise in the construction. All assumptions correspond to explicit statements in the HKDS specification or to standard cryptographic models for Keccak based primitives. No symbolic domain tags or idealized labels are introduced; domain separation follows from structural distinctions in the inputs to SHAKE and KMAC.

#### 4.1 Cryptographic Primitive Assumptions

HKDS relies on two Keccak based primitives: SHAKEp and KMACp. Both are modeled in this analysis using standard XOF based assumptions.

##### SHAKEp

SHAKEp is modeled as an extendable output pseudorandom function on variable length inputs. For any fixed rate  $\rho$ , SHAKEp produces an output stream indistinguishable from a uniform bitstring to any efficient adversary that does not know the input. Its internal state evolution follows the Keccak permutation as defined in FIPS 202 and SP 800 185.

##### KMACp

KMAC is modeled as a pseudorandom function in its key input. For any fixed rate  $\rho$ , KMACp( $K, M$ ) is indistinguishable from a random function indexed by  $K$ . The function inherits the domain separation guarantees of Keccak's cSHAKE based customization mechanism.

These assumptions are standard for Keccak derived constructions and are consistent with the HKDS specification.

#### 4.2 Structural Domain Separation in HKDS

HKDS does not rely on explicit literal domain strings for domain separation. Instead, it achieves separation through the structural arrangement of inputs to SHAKE and KMAC. The following inputs are all distinct in purpose, length, and ordering, and no two stages of the protocol feed the same structured input into the same primitive.

### **Embedded Device Key**

The embedded device key is

$$\text{EDK} = \text{SHAKE}_\rho(\text{DID} \parallel \text{BDK}).$$

The concatenation  $\text{DID} \parallel \text{BDK}$  is unique to device key derivation. No other HKDS component uses this input structure.

### **Token Derivation**

Tokens are derived exclusively by the server as

$$\text{TOK} = \text{SHAKE}_\rho(\text{CTOK} \parallel \text{STK}),$$

where  $\text{CTOK} = \text{TRC} \parallel \text{FN} \parallel \text{DID}$ .

The tuple  $(\text{TRC}, \text{FN}, \text{DID})$  ensures that each  $\text{CTOK}$  value is bound to the device identity and the current transaction region. This structure is not reused elsewhere.

### **Token Encryption Pad**

The encryption pad used to protect  $\text{TOK}$  in transit is

$$\text{Kpad} = \text{SHAKE}_\rho(\text{CTOK} \parallel \text{EDK}).$$

Although  $\text{CTOK}$  appears again, the presence of  $\text{EDK}$  instead of  $\text{STK}$  is a structural difference that enforces separation from both the token derivation and the device key derivation stages.

### **Transaction Key Cache**

The transaction key cache depends jointly on  $\text{TOK}$  and  $\text{EDK}$ :

$$\text{TKC} = \text{SHAKE}_\rho(\text{TOK} \parallel \text{EDK}).$$

The concatenation order  $\text{TOK} \parallel \text{EDK}$  is unique to cache generation. No other input uses this arrangement or ordering.

### **Message Authentication**

KMAC uses the transaction key provided directly from TKC. The token authentication uses

$$\text{TAGt} = \text{KMACp(EDK, ETOK, custom = TMS)},$$

where  $\text{TMS} = \text{KSN} \parallel \text{MACname}$ .

The customization string TMS is unique to the token MAC and binds the tag to the device identity, key identifier, and transaction counter.

These structural distinctions avoid cross protocol collisions and ensure that no single value is reused as a key or seed in more than one derivation path.

#### 4.3 Erasure Assumptions and State Exposure

HKDS specifies that active token and cache material must be erased once consumed. The forward secrecy claims rely on the assumption that, after cache and token erasure, compromise of the remaining client state does not reveal previously used transaction keys.

The security model assumes the following:

1. **Tokens and cache entries are erased promptly after use.**

The adversary cannot recover past TK values once both the token and the corresponding cache segment have been erased.

2. **The adversary may later compromise EDK or other long term client state.**

However, because recovery of TOK requires access to both ETOK and CTOK, and because ETOK is not retained after cache erasure, past keys remain infeasible to reconstruct.

3. **Recorded network traffic is available to the adversary.**

The adversary may retain ETOK, CTOK, KSN, and all ciphertexts, but cannot recover past TK values without the corresponding token, which has already been erased.

These are explicit operational requirements and must be satisfied by any secure implementation.

#### 4.4 Adversarial Capabilities

The adversary is modeled as having:

1. full access to all public protocol messages, including KSN, CTOK dependent fields, ETOK, ciphertexts, and tags
2. the ability to delay, replay, or reorder messages
3. no access to BDK or STK
4. no physical compromise of the MDK or secure server environment
5. possible later compromise of client state after token erasure, as permitted by the forward secrecy model

The adversary cannot extract keys from Keccak's internal state, nor can it cause the primitives to deviate from their specification.

#### 4.5 Summary of Security Model

HKDS relies on three primary classes of assumptions:

1. **Keccak XOF security**, guaranteeing pseudo-randomness of SHAKE and KMAC outputs.
2. **Structural domain separation**, achieved through distinct concatenation structures at each derivation stage.
3. **Timed erasure of sensitive material**, ensuring that past transaction keys cannot be recovered after the token and its cache segment have been erased.

These assumptions form the foundation for the proofs developed in later chapters.

#### 4.6 Predictive resistance

Predictive resistance requires that an adversary who learns all material associated with the current token epoch cannot compute transaction keys from a future epoch without access to the next server token.

Formally, consider an adversary that knows:

1. the current embedded device key EDK
2. the current token TOK and all transaction keys  $K_1, \dots, K_t$  derived from  $TKC = \text{SHAKE}_p(TOK \parallel EDK)$

3. the full transcript of all protocol messages, including KSN and CTOK values for the current and prior epochs.

After the current token has been fully consumed and erased, the server issues a new token  $TOK' = \text{SHAKE}_p(CTOK' \parallel STK)$  for the same device. The next cache is  $TKC' = \text{SHAKE}_p(TOK' \parallel EDK)$ , with keys  $K_1', K_2'$ , and so on.

Since STK is server confined and  $TOK'$  depends on both  $CTOK'$  and STK, the adversary cannot compute  $TOK'$  from public data or from any client state. Under the assumption that  $\text{SHAKE}_p$  behaves as a pseudorandom function on  $(CTOK' \parallel STK)$ , the distribution of  $TOK'$  is independent of previous tokens, even if the adversary knows all prior  $CTOK$  values.

Consequently, future transaction keys  $K_i'$  are pseudorandom from the adversary's point of view, and no efficient strategy can predict them from EDK, prior tokens, or prior cache outputs. Predictive resistance therefore holds across token epochs, provided that STK remains secret and that CTOK values are not reused for the same device.

#### **4.7 Authenticity and integrity**

HKDS provides authenticity and integrity at two levels: individual messages and token transport.

For message protection, each transaction uses two consecutive cache entries. The first entry  $K_{enc}$  is used to encrypt the message by XOR, and the second entry  $K_{mc}$  is used as the key to KMAC. Given ciphertext  $C$  and optional associated data  $A$ , the tag is

$$\text{TAG} = \text{KMAC}_p(K_{mc}, C \parallel A, \text{custom} = \text{empty}).$$

Under the assumption that  $\text{KMAC}_p$  is a pseudorandom function in its key input, and that each  $K_{mc}$  is used at most once, the advantage of any adversary in forging a valid tag without knowledge of  $K_{mc}$  is bounded by approximately  $2^{-n}$  for output length  $n$ . Since  $K_{mc}$  is never reused, the adversary has at most one meaningful attempt per key, so unforgeability is preserved up to the birthday bound on the tag length.

For token transport, the server authenticates the encrypted token  $ETOK$  with a tag

$$\text{TAG}_t = \text{KMAC}_p(EDK, ETOK, \text{custom} = TMS),$$

where  $TMS = KSN \parallel \text{MACname}$ . The client recomputes  $TMS$  from its  $KSN$  and the agreed MAC name, then verifies  $\text{TAG}_t$  using the same  $EDK$ . This binds  $ETOK$  to the specific

device identity, key identifier, and transaction counter encoded in KSN. Under the KMAC PRF assumption and secrecy of EDK, any successful forgery of TAGt provides a distinguishing attack on KMAC. The probability of undetected modification of ETOK is therefore negligible.

#### 4.8 Determinism and recoverability

The construction of HKDS is fully deterministic. For any fixed parameters (BDK, STK, KID, DID, KTC), the pair of client and server computations must produce identical transaction keys.

On the client side, the effective inputs are the provisioned EDK, the recovered token TOK, and the transaction counter KTC. The client computes  $TKC = \text{SHAKE}_\rho(TOK \parallel EDK)$  and selects the appropriate entries from TKC based on KTC.

On the server side, given (BDK, STK, KID, DID, KTC) extracted from the MDK and KSN, the server performs:

1.  $EDK = \text{SHAKE}_\rho(DID \parallel BDK)$
2.  $CTOK = TRC \parallel FN \parallel DID$ , where TRC is derived from KTC and the cache size
3.  $TOK = \text{SHAKE}_\rho(CTOK \parallel STK)$
4.  $TKC = \text{SHAKE}_\rho(TOK \parallel EDK)$

The server then selects the same logical positions in TKC as the client, using the same counter KTC and cache parameters. Because each step is a deterministic function of its inputs, and the inputs coincide when KSN and counters are synchronized, the resulting transaction keys are identical.

Recoverability therefore holds for all transactions so long as the client and server agree on KID, DID, and KTC and the MDK components (BDK, STK) remain unchanged.

#### 4.9 Domain separation proof sketch

Domain separation in HKDS follows from the fact that each derivation stage feeds a structurally distinct input into SHAKE or KMAC. There is no stage at which the same primitive is applied to two different conceptual tasks with identical input structure.

The relevant inputs are:

1. Device key:  $DID \parallel BDK$ , used only in  $EDK = \text{SHAKE}_\rho(DID \parallel BDK)$ .

2. Token seed:  $CTOK \parallel STK$ , used only in  $TOK = \text{SHAKE}_\rho(CTOK \parallel STK)$ .
3. Token pad:  $CTOK \parallel EDK$ , used only in  $Kpad = \text{SHAKE}_\rho(CTOK \parallel EDK)$ .
4. Cache seed:  $TOK \parallel EDK$ , used only in  $TKC = \text{SHAKE}_\rho(TOK \parallel EDK)$ .
5. Token MAC: key  $EDK$  with message  $ETOK$  and customization  $TMS$ .
6. Message MAC: key  $Kmc$  with message  $C \parallel A$  and empty customization.

Under the  $\text{SHAKE}_\rho$  pseudo-randomness assumption, any efficient adversary that distinguishes outputs at one derivation stage from random must either invert or distinguish  $\text{SHAKE}_\rho$  on the corresponding input structure. Since no other stage uses the same structure, such an adversary cannot transfer that advantage to another stage. Similarly, under the KMAC PRF assumption, tags computed with  $EDK$  and tags computed with  $Kmc$  are independent, because they use different keys and different message formats.

In particular:

- Knowledge of  $EDK$  and  $CTOK$  does not reveal  $STK$ , since recovering  $STK$  from  $TOK = \text{SHAKE}_\rho(CTOK \parallel STK)$  is equivalent to inverting  $\text{SHAKE}_\rho$ .
- Knowledge of  $TOK$  and  $DID$  does not reveal  $BDK$ , since recovering  $BDK$  from  $EDK = \text{SHAKE}_\rho(DID \parallel BDK)$  is also equivalent to inverting  $\text{SHAKE}_\rho$ .
- Knowledge of transaction keys  $K_i$  derived from  $TKC$  does not reveal  $TOK$  or  $EDK$ , since this again requires inverting  $\text{SHAKE}_\rho$  on  $TOK \parallel EDK$ .

These observations show that, given the primitive assumptions, the different derivation domains are isolated and there is no feasible way to reuse information from one stage to attack another.

#### 4.10 Summary of properties

The corrected model of HKDS yields the following properties under the assumed security of  $\text{SHAKE}_\rho$  and  $\text{KMAC}_\rho$  and under the operational constraints specified in Chapter 2.

First, each transaction key  $K_i$  in the cache  $TKC$  is pseudorandom from the adversary's perspective, provided that  $TOK$  and  $EDK$  remain unknown. The XOR encryption  $C = M \oplus$

$K_{enc}$  is therefore indistinguishable from a one-time pad encryption for single use of  $K_{enc}$ .

Second, predictive resistance holds across token epochs. Even if an adversary learns EDK, all transaction keys in the current cache, and all protocol transcripts, future tokens remain hidden because they depend on the server confined key STK and fresh CTOK values. The next cache TKC' is seeded with a new TOK' and is therefore independent of previous caches.

Third, authenticity and integrity are provided by KMACp both at the message level and for token transport. Since  $K_{mc}$  and EDK are never reused across unrelated tags, and KMAC is modeled as a PRF, the probability of successful forgery without the corresponding key is negligible.

Fourth, determinism and recoverability are guaranteed by the structure of the derivation chain. Given synchronized counters and consistent MDK and KSN values, client and server compute identical transaction keys without requiring persistent per device state on the server.

Finally, domain separation is enforced by the distinct input structures supplied to SHAKEp and KMACp at each derivation stage. No single key or seed value is reused in more than one role with the same primitive and input format, which prevents unintended key overlap between device keys, tokens, caches, and authentication tags.

These properties provide the foundation for the more detailed cryptanalytic evaluation developed in the subsequent chapters.

## Chapter 5: Cryptanalysis and Robustness Evaluation

### Cryptanalytic Evaluation

This chapter evaluates the HKDS construction using the corrected derivation model and the security assumptions stated in Chapters 2 through 4. The analysis proceeds in terms of the actual functions used by HKDS: SHAKEp for key derivation and expansion, and KMACp for authentication. All conclusions are based on the behavior of the real design and not on idealized variants.

The evaluation covers the principal cryptanalytic dimensions: resistance to key compromise, stability under state exposure, forward secrecy across token epochs, predictive resistance, recoverability, and robustness against tampering or replay. Each property follows from the derivation equations presented in Chapter 3 and the assumptions summarized in Chapter 4.

## 5.1 Derivation-chain exposure

The derivation chain of HKDS consists of:

1.  $\text{EDK} = \text{SHAKE}_\rho(\text{DID} \parallel \text{BDK})$
2.  $\text{TOK} = \text{SHAKE}_\rho(\text{CTOK} \parallel \text{STK})$
3.  $\text{TKC} = \text{SHAKE}_\rho(\text{TOK} \parallel \text{EDK})$
4.  $\text{K}_{\text{enc}}$  and  $\text{K}_{\text{mc}}$  extracted from  $\text{TKC}$
5.  $\text{TAG} = \text{KMAC}_p(\text{K}_{\text{mc}}, \text{C} \parallel \text{A})$
6.  $\text{TAGt} = \text{KMAC}_p(\text{EDK}, \text{ETOK}, \text{custom} = \text{TMS})$

An adversary who compromises any single output of this chain cannot feasibly invert  $\text{SHAKE}_\rho$  to recover any preceding input. For instance, recovering  $\text{BDK}$  from  $\text{EDK}$  requires inverting  $\text{SHAKE}_\rho$  on  $\text{DID} \parallel \text{BDK}$ , and recovering  $\text{STK}$  from  $\text{TOK}$  requires inverting  $\text{SHAKE}_\rho$  on  $\text{CTOK} \parallel \text{STK}$ . These operations are assumed infeasible.

Similarly, learning any subset of transaction keys from  $\text{TKC}$  does not reveal  $\text{TOK}$  or  $\text{EDK}$ , since the mapping  $(\text{TOK} \parallel \text{EDK}) \rightarrow \text{TKC}$  is a pseudo-random function on its input. The adversary therefore gains no advantage in attacking earlier derivation stages by observing or compromising later ones.

## 5.2 Token secrecy and token-transport integrity

The server transmits an encrypted token  $\text{ETOK}$  constructed as

$$\text{ETOK} = \text{TOK} \oplus \text{SHAKE}_\rho(\text{CTOK} \parallel \text{EDK}).$$

Under the  $\text{SHAKE}_\rho$  pseudo-randomness assumption, the pad  $\text{SHAKE}_\rho(\text{CTOK} \parallel \text{EDK})$  is indistinguishable from a uniform string even to an adversary who knows  $\text{CTOK}$  and  $\text{DID}$ . Since  $\text{EDK}$  is unknown to the adversary, the token ciphertext reveals no information about  $\text{TOK}$ .

Integrity of the token is provided by

$\text{TAGt} = \text{KMACp}(\text{EDK}, \text{ETOK}, \text{custom} = \text{TMS}),$

where TMS includes KSN and the MAC name agreed by both sides. Any tampering with ETOK results in an invalid tag unless the adversary can forge KMACp under the key EDK, which contradicts the PRF assumption. Since EDK does not leave the device except implicitly in derived outputs, token transport provides both confidentiality and authenticity.

### 5.3 Message confidentiality

HKDS encrypts each message block by

$$C = M \oplus K_{\text{enc}},$$

where  $K_{\text{enc}}$  is the next unused transaction key from TKC. Each  $K_{\text{enc}}$  is a block of output from  $\text{SHAKE}_p(\text{TOK} \parallel \text{EDK})$ . Under the XOF pseudo-randomness assumption,  $K_{\text{enc}}$  is computationally indistinguishable from a uniform bitstring, and since it is used once, the resulting encryption is equivalent to a one-time pad on that block.

Message confidentiality is therefore reduced to the secrecy of TOK and EDK. Without knowledge of both inputs to the  $\text{SHAKE}_p$  instance that produced  $K_{\text{enc}}$ , no adversary can distinguish  $C$  from a uniformly random string.

### 5.4 Message authentication

Integrity of ciphertexts and optional associated data is provided by

$$\text{TAG} = \text{KMACp}(K_{\text{mc}}, C \parallel A),$$

where  $K_{\text{mc}}$  is the next unused cache entry. Since  $K_{\text{mc}}$  is never reused and KMAC is a PRF in its key input, the probability that an adversary can produce a valid tag without knowledge of  $K_{\text{mc}}$  is bounded by approximately  $2^{-n}$  for output length  $n$ .

Authenticated encryption in HKDS is therefore secure in both the unforgeability and chosen ciphertext senses, provided that  $K_{\text{mc}}$  remains secret and that no transaction key is reused.

### 5.5 Forward secrecy across token epochs

Forward secrecy in HKDS is point-in-time and depends on erasure of token and cache material after use. Let TOK and TKC represent the state of an earlier epoch. Once TOK

and TKC have been erased, later compromise of the client that reveals only EDK does not permit the adversary to reconstruct earlier transaction keys.

This is because ETOK for earlier epochs is not retained on the client, and reconstruction of TOK requires knowledge of both ETOK and the pad  $\text{SHAKE}_\rho(\text{CTOK} \parallel \text{EDK})$ .

If ETOK and the corresponding pad are no longer present, the adversary cannot recover TOK. Without TOK, the adversary cannot regenerate TKC, since  $\text{TKC} = \text{SHAKE}_\rho(\text{TOK} \parallel \text{EDK})$ .

Forward secrecy in HKDS is therefore conditional on correct erasure practices but holds under that operational assumption.

## 5.6 Predictive resistance in future epochs

Predictive resistance requires that the adversary cannot compute future transaction keys even if it knows all material from previous epochs. The next epoch uses

$$\text{TOK}' = \text{SHAKE}_\rho(\text{CTOK}' \parallel \text{STK})$$

and

$$\text{TKC}' = \text{SHAKE}_\rho(\text{TOK}' \parallel \text{EDK}).$$

Since CTOK' differs from CTOK by inclusion of an updated transaction region code, and since STK is never exposed outside the server, TOK' is pseudorandom and independent of TOK. Without TOK', the adversary cannot compute TKC' or any future Kenc or Kmc values.

Predictive resistance thus follows directly from the secrecy of STK and non-reuse of CTOK values.

## 5.7 Replay and substitution

Replay or substitution of ETOK or of message ciphertexts is detected by HKDS through two independent mechanisms.

First, each token epoch is bound to a distinct CTOK value. A token replayed outside its intended region results in recovery of an incorrect cache or a mismatch between expected and actual counter indices.

Second, each ciphertext includes a KMAC tag computed under a unique Kmc. Replayed messages are rejected because the server transaction counter will not match the counter

encoded in KSN, and substituted messages fail verification unless the adversary can forge KMAC tags.

Replay resistance is therefore inherent to both the token structure and the counter synchronized cache.

### **5.8 State rollback and counter recovery**

HKDS requires monotonic advancement of the transaction counter KTC. If a client attempts to reuse an older KTC, the server detects a mismatch when validating KSN and rejects the request. Conversely, if a client receives responses out of order or loses synchronization, it must discard its current cache and request a new token from the server, ensuring that both sides realign their internal states.

The deterministic nature of the derivation chain guarantees that any counter value KTC corresponds to a unique set of derived keys, making rollback attacks ineffective so long as counter reuse is prevented.

### **5.9 Robustness against subkey reuse**

Subkey reuse would weaken both confidentiality and integrity. HKDS avoids reuse by consuming each transaction key exactly once and requiring clients to request new tokens whenever the cache is exhausted. Since TKC is derived from TOK and EDK and all CTOK values are unique for each region, each token epoch produces a distinct TKC, and each Kenc and Kmc pair is used for only one operation.

The cryptanalytic strength therefore depends on correct counter handling and timely token refresh, both of which are enforced by the protocol structure.

### **5.10 Summary**

HKDS provides confidentiality, authenticity, forward secrecy across token epochs, predictive resistance, and deterministic recoverability under the assumption that SHAKEp and KMACp behave as ideal pseudorandom functions and that token and cache erasure is performed correctly. The design's structural domain separation ensures that no derivation stage compromises another, and the synchronized counter mechanism prevents replay and subkey reuse. Under these conditions, HKDS meets its stated cryptographic objectives.

## Chapter 6: Implementation and Performance Evaluation

### 6.1 Computational structure

The HKDS implementation relies exclusively on Keccak-based primitives. All key derivations, token computations, and message protections are realized through the same sponge function with rate and capacity parameters dependent on the security strength. No external cipher, key schedule, or block transformation is employed. This simplifies both software and hardware implementations and provides a uniform execution profile across all operational phases.

Each token or message operation involves the following computational cost:

- One invocation of SHAKE for key derivation (EDK, token, or cache).
- One invocation of KMAC for authentication.
- Simple XOR operations for encryption and decryption.

The overall computational complexity per transaction is dominated by Keccak permutations. For SHAKE256, this corresponds to 24 rounds per 136-byte block, yielding predictable timing behavior.

### 6.2 Memory and state requirements

The HKDS client maintains:

1. A 32-byte EDK.
2. The current token T (32 or 64 bytes, depending on the mode).
3. The current cache buffer, typically a few hundred bytes of pre-derived keys.
4. A monotonic 32-bit counter KTC.

The server retains:

- The BDK and STK roots (each one secret block).
- Persistent device metadata (DID, KID, KTC).

The combined memory footprint is minimal, well within the capabilities of low-power microcontrollers and secure elements used in payment terminals.

### 6.3 Synchronization and throughput

The deterministic derivation chain ensures that the server can regenerate any transaction key without transmitting additional state. Synchronization relies only on KSN, eliminating multi-round negotiation or handshake procedures.

A token refresh involves a single request-response pair and is required only after cache exhaustion. For example, if a cache contains 1000 transaction keys, 1000 secure operations can occur before another token request. This amortizes the network cost of key distribution and allows large-scale systems to maintain continuous operation with minimal bandwidth overhead.

Throughput is therefore bounded primarily by the performance of SHAKE. On modern embedded processors, a single 24-round Keccak permutation executes in microseconds, yielding transaction speeds suitable for real-time payment processing.

#### **6.4 Determinism and reproducibility**

All key derivations are strictly deterministic functions of (EDK, T, KTC). No runtime entropy source influences the outcome. This guarantees reproducibility across client and server and permits offline verification or forensic reconstruction of any transaction key by recomputing the derivation chain.

This property also simplifies compliance verification, as all outputs are auditably tied to known cryptographic roots.

#### **6.5 Error handling and recovery behavior**

When synchronization errors occur, such as cache exhaustion or counter mismatch, the implementation discards the current cache and requests a new token. The server recomputes from the stored KTC to ensure continuity. Because all functions are idempotent with respect to inputs, repeated computation yields identical keys, preventing drift between devices and server.

Cache regeneration and token refresh are lightweight operations involving only SHAKE and KMAC computations. There is no dependency on database locks or inter-device communication, allowing recovery even under high load or intermittent connectivity.

#### **6.6 Implementation security considerations**

##### **6.6.1 Timing uniformity**

The Keccak permutation uses fixed iteration counts and uniform memory access, eliminating timing leaks related to key-dependent branching. The same function paths are executed regardless of message length or key content.

### **6.6.2 Memory zeroization**

The reference code clears intermediate buffers after finalization, ensuring that transient states (e.g., partial sponge buffers) do not persist in memory. This reduces exposure to local memory inspection attacks.

### **6.6.3 Key storage**

The specification requires that BDK, STK, and EDK values be stored within tamper-resistant hardware. If this requirement is met, physical compromise is limited to side-channel attacks on SHAKE, which remain infeasible given constant-time implementation and noise-dominated power signatures.

### **6.6.4 Parallelism**

Although HKDS is primarily sequential, the cache-generation phase can be parallelized by producing multiple SHAKE streams from the same seed offset. This optimization enables higher throughput on multicore systems without altering the deterministic mapping between indices and keys.

## **6.7 Interoperability**

HKDS is compatible with standard Keccak libraries that implement FIPS 202 and SP 800-185. It does not depend on platform-specific instructions or non-standard state layouts. The same derivation routines execute identically on both little-endian and big-endian systems due to explicit byte-ordering controls in the provided code.

As a result, interoperability between heterogeneous systems; such as embedded clients, servers, and hardware security modules, is maintained without normalization steps or encoding conversions.

## **6.8 Scalability and performance summary**

The deterministic, server-stateless design permits an arbitrary number of clients to operate concurrently without additional per-session storage. Each transaction key can be recomputed in constant time, and token issuance scales linearly with the number of connected devices.

The total computational cost for each secure transaction is limited to:

- One SHAKE call for cache expansion (precomputed or on-demand).
- One KMAC verification or generation.

Given these properties, HKDS achieves high throughput and low latency while maintaining strict forward secrecy and recoverability.

## **Chapter 7: Security Model Evaluation and Comparative Discussion**

### **7.1 Overview**

This chapter evaluates the HKDS design under established cryptographic security models and compares its resilience and structural properties with those of related key-distribution schemes such as DUKPT and conventional symmetric derivation hierarchies used in payment systems. The analysis considers both theoretical soundness and operational durability, integrating the formal properties defined in earlier chapters with the behavioral outcomes of the provided implementation.

### **7.2 Evaluation under standard models**

#### **7.2.1 Indistinguishability under chosen-plaintext attack (IND-CPA)**

Encryption in HKDS uses a stream derived from SHAKE keyed by the pair (TOK, EDK); each Kenc is a block of this stream, and encryption is performed as  $C = M \oplus K_{enc}$ .

Because each Kenc is a one-time key obtained from a pseudorandom stream seeded with a unique token and device key, ciphertexts corresponding to distinct messages are computationally indistinguishable from random bitstrings. Given that SHAKE behaves as a pseudorandom function with uniform output distribution, the adversary's advantage in distinguishing ciphertexts is negligible in  $2^{-n}$ .

#### **7.2.2 Unforgeability under chosen-message attack (UF-CMA)**

The KMAC-based authentication mechanism provides message integrity and authenticity under the assumption that KMAC functions as a secure MAC. Since each KMAC key ( $K_{mc}$ ) is unique and never reused, the standard security bound for one-time MACs applies, producing a forgery probability of  $2^{-n}$ . Even with  $q$  attempted forgeries, the total advantage remains below  $q/2^n$ , which is negligible for realistic system scales.

### 7.2.3 Key-compromise impersonation (KCI)

In HKDS, compromise of EDK on a single client does not enable impersonation of other devices because each device's derivation path is bound to its DID. Similarly, compromise of STK does not enable impersonation of any client without EDK. The two-key model prevents cross-domain forgery; adversaries require both to reconstruct valid tokens and message keys. This division resists typical KCI attacks observed in single-root DUKPT derivations.

### 7.2.4 Replay and desynchronization resilience

The KTC counter provides temporal ordering for all transactions. Messages with repeated or stale KTC values are rejected, rendering replay attacks infeasible. The deterministic server recomputation model ensures that loss of synchronization can be corrected without breaking forward secrecy, satisfying robustness against desynchronization; a property not guaranteed in traditional DUKPT.

## 7.3 Comparison to DUKPT

DUKPT uses a tree-structured key derivation model where each transaction key is deterministically derived from a single base key through a series of non-invertible operations involving variant bits and registers. This model enforces one-time key use but introduces heavy key-management complexity and limited scalability, as each device must independently derive all keys from its initial key seed.

HKDS replaces the tree expansion with a two-key system based on SHAKE, which yields several measurable advantages:

- **Entropy refresh:** Tokens issued by the server introduce new entropy periodically, ensuring forward secrecy without rekeying the device. DUKPT's deterministic expansion cannot achieve this without replacing the base key.
- **Scalability:** Server statelessness and deterministic recomputation eliminate the need for per-device key trees or persistent key tables.
- **Authenticity:** KMAC adds authenticated integrity; DUKPT relies on implicit trust without built-in message authentication.

- **Cryptographic agility:** HKDS supports 128–512-bit configurations and direct integration with Keccak-based primitives approved under FIPS 202, while DUKPT remains limited to legacy symmetric ciphers such as 3DES or AES-128.

In cryptanalytic terms, HKDS achieves equivalent or superior strength while simplifying operational overhead.

#### **7.4 Comparison to generic KDF-based architectures**

Conventional key hierarchies often apply HKDF or PBKDF constructions to derive multiple keys from a master secret. These typically lack counter-based synchronization and require the client and server to store per-session nonces or salts. HKDS avoids this by embedding the counter in the KSN and using deterministic derivations, which produce identical results on both sides without additional state exchange.

Unlike HKDF, which depends on HMAC and therefore on iterative SHA-2 hashing, HKDS inherits the diffusion and uniform output characteristics of Keccak’s sponge function, offering improved throughput and simpler domain separation. The resulting design eliminates feedback loops and salt storage while maintaining equivalent resistance to key-collision and preimage attacks.

#### **7.5 Forward secrecy and token rotation policy**

Forward secrecy in HKDS is defined per token epoch. As each new token introduces fresh entropy derived from the server’s STK, keys from prior epochs become computationally independent. The security of this property depends on the rotation interval: shorter token lifetimes provide stronger containment against device compromise.

The mathematical analysis confirms that even with full exposure of one token epoch, subsequent tokens remain unpredictable without STK. This periodic renewal mechanism aligns HKDS with best practices in symmetric key ratcheting and provides a degree of dynamic forward secrecy uncommon in symmetric-only designs.

#### **7.6 Robustness under compromised subsystems**

##### **7.6.1 Compromised device**

If an attacker captures EDK and current cache state, they can decrypt transactions within the active token epoch but cannot generate valid tags or new cache entries once the

server issues a new token. Because the server uses STK to generate tokens, it can revoke compromised devices simply by discontinuing token service or rotating STK.

### 7.6.2 Compromised server

Exposure of STK allows regeneration of tokens but not decryption of stored transactions, as EDK remains device-specific. Without EDK, ciphertexts remain cryptographically opaque. The dual-secret model thus prevents complete system compromise through a single breach.

### 7.6.3 Simultaneous exposure

If both BDK and STK are lost, all derived keys become recoverable. This represents the root compromise case. The design's reliance on hardware isolation and strict key separation mitigates this scenario by requiring dual authorization for access, consistent with dual-control policies used in financial HSM deployments.

## 7.7 Structural advantages

The HKDS hierarchy achieves three significant design improvements over conventional schemes:

1. **Complete determinism with bounded recomputation:** Every key can be recomputed by the server using minimal state, enabling forensic traceability without exposing keys in transit.
2. **Strict key separation through structural domain separation:** Each derivation layer uses distinct input structures, preventing cross-domain key collisions or misuse.
3. **Symmetric forward secrecy without asymmetric primitives:** HKDS accomplishes forward secrecy using only Keccak-based symmetric operations, simplifying compliance for constrained environments where public-key computation is impractical.

## 7.8 Summary

The HKDS model satisfies the principal goals of confidentiality, integrity, and recoverability within an ideal-primitive framework. The separation of derivation roots ensures that partial compromise does not cascade into universal failure. Its deterministic

nature simplifies audit and recovery, while SHAKE and KMAC provide modern cryptographic strength consistent with current NIST standards.

In comparison to legacy symmetric key distribution schemes, HKDS offers a clearer separation of roles, stronger integrity protection, and better scalability without introducing new theoretical weaknesses.

## Chapter 8: Formal Security Proofs and Reduction Analysis

### Formal Security Analysis

This chapter presents formal analyses of the principal security properties claimed for HKDS. The treatment follows the real construction as defined in Chapter 3 and the adversarial model specified in Chapter 2. The analysis uses only the actual functions present in HKDS: SHAKE $\rho$  for key derivation and expansion, and KMAC $\rho$  for authentication. All proofs are carried out under the idealized assumptions stated in Chapter 4 that SHAKE $\rho$  behaves as a pseudorandom function on its input and that KMAC $\rho$  behaves as a pseudorandom function in its key input. No idealized domain identifiers or auxiliary keys are introduced, and all derivation paths use their correct input concatenation forms.

The properties examined in this chapter are message confidentiality, message authenticity, token authenticity, forward secrecy across token epochs, predictive resistance, determinism and recoverability, and domain separation.

#### 8.1 Message confidentiality

Let  $K_{enc}$  be the next unused transaction key in TKC, where

$$TKC = \text{SHAKE}_\rho(TOK \parallel EDK)$$

and

$$TOK = \text{SHAKE}_\rho(CTOK \parallel STK).$$

The client encrypts each message block by

$$C = M \oplus K_{enc}.$$

Since  $K_{enc}$  is a substring of  $\text{SHAKE}_\rho(TOK \parallel EDK)$ , and since  $\text{SHAKE}_\rho$  is modeled as a pseudorandom function on fixed inputs, the distribution of  $K_{enc}$  is computationally

indistinguishable from a random bitstring of the same length. Provided that  $K_{enc}$  is used once, the ciphertext  $C$  is indistinguishable from a one-time pad encryption.

### **Indistinguishability experiment.**

In an IND-CPA experiment, an adversary chooses two messages  $M_0$  and  $M_1$  of equal length. The challenger samples a bit  $b$ , computes  $C = M_b \oplus K_{enc}$ , and provides  $C$ . Because  $K_{enc}$  is indistinguishable from a uniform string and is independent of  $M_b$ , no efficient adversary can distinguish which message was encrypted with non-negligible advantage.

### **Conclusion.**

Message confidentiality in HKDS reduces directly to the pseudo-randomness of SHAKEp. The encryption scheme is therefore IND-CPA secure under the stated assumptions.

## **8.2 Message authenticity**

HKDS authenticates ciphertexts and associated data using

$$TAG = \text{KMACp}(K_{mc}, C \parallel A),$$

where  $K_{mc}$  is the next unused cache entry and has not been used in any previous tag computation.

Since  $K_{mc}$  is secret and is used exactly once, and since KMAC is modeled as a pseudo-random function in its key input, the adversary's ability to produce a valid tag without knowledge of  $K_{mc}$  reduces to the advantage of distinguishing KMAC from a random function.

### **Unforgeability experiment.**

An adversary obtains tags for messages of its choice, each under distinct keys  $K_{mc}$  that it does not possess. To yield a successful forgery, the adversary must produce  $(C, A, TAG)$  such that  $TAG$  equals  $\text{KMACp}(K_{mc}, C \parallel A)$  for some unused  $K_{mc}$ . Under the PRF assumption, this success probability is bounded by approximately  $2^{-n}$  for  $n$ -bit tags.

### **Conclusion.**

HKDS message authentication is UF-CMA secure under the KMAC PRF assumption and single use of each  $K_{mc}$  value.

## **8.3 Token authenticity**

The server authenticates  $E_{TOK}$  using

$TAG_t = KMAC_p(EDK, ETOK, \text{custom} = TMS)$ ,

where  $TMS = KSN \parallel \text{MACname}$ .

### Security rationale.

The adversary does not know EDK. Any modification of ETOK or substitution of a forged ETOK' causes verification to fail unless the adversary can produce a valid TAGt' under the same TMS. Since KMACp is a pseudorandom function in its key input, any distinguisher capable of forging TAGt induces a PRF distinguisher on KMAC.

### Conclusion.

Token authenticity reduces to the PRF security of KMACp under the key EDK. The probability of undetected token modification is negligible.

## 8.4 Forward secrecy across token epochs

Forward secrecy in HKDS relies on the erasure of token and cache material once consumed. Let TOK and TKC represent the state of an earlier epoch. Once the client has erased TOK, Kpad, and all Kenc and Kmc values derived from that epoch, later compromise of the client state reveals only EDK and public transcript data.

The adversary can retain CTOK and ETOK from previous epochs, but without access to the pad

$Kpad = \text{SHAKE}_p(CTOK \parallel EDK)$ ,

which is erased after use, the adversary cannot recover TOK. Without TOK, the adversary cannot regenerate TKC, since  $TKC = \text{SHAKE}_p(TOK \parallel EDK)$ .

### Conclusion.

Forward secrecy holds for previous epochs as long as token and cache erasure is performed correctly and EDK is not combined with preserved ETOK and CTOK from the same epoch. This matches the forward secrecy model defined in the HKDS specification.

## 8.5 Predictive resistance

Predictive resistance requires that the adversary cannot compute future transaction keys from previous epochs, even if all material from earlier epochs is known. Future epochs use tokens of the form

$TOK' = \text{SHAKE}_p(CTOK' \parallel STK)$ .

Since CTOK' differs from CTOK by the transaction region code and since STK is server confined, an adversary cannot derive TOK' without knowledge of STK. Because TKC' = SHAKEp(TOK' || EDK), and because TOK' is pseudo-random, no adversary can compute future Kenc or Kmc values without the new token.

### **Conclusion.**

Predictive resistance holds across token epochs under secrecy of STK and uniqueness of CTOK values.

## **8.6 Determinism and recoverability**

HKDS requires client and server to produce identical Kenc and Kmc values when operating with matching DID, KID, BDK, STK, and KTC. The derivation sequence

$$\text{EDK} = \text{SHAKEp}(\text{DID} \parallel \text{BDK})$$

$$\text{TOK} = \text{SHAKEp}(\text{CTOK} \parallel \text{STK})$$

$$\text{TKC} = \text{SHAKEp}(\text{TOK} \parallel \text{EDK})$$

is deterministic. For identical inputs, both sides obtain the same EDK, the same TOK, and the same TKC. Recoverability follows when KTC is synchronized and the server reconstructs all inputs from KSN and MDK.

### **Conclusion.**

Determinism and recoverability hold unconditionally for all synchronized transactions and follow directly from the determinism of SHAKEp.

## **8.7 Domain separation**

HKDS does not use symbolic domain tags. Instead, domain separation follows from the structural form of the inputs fed into SHAKEp and KMACp:

1. EDK uses DID || BDK, unique to device key derivation.
2. TOK uses CTOK || STK, unique to token derivation.
3. Kpad uses CTOK || EDK, unique to token encryption.
4. TKC uses TOK || EDK, unique to cache generation.
5. KMACp for ETOK uses key EDK and customization TMS.
6. KMACp for ciphertext uses Kmc and empty customization.

No two derivation steps feed the same structure into the same primitive. Any adversary distinguishing one derivation stage from random cannot translate that advantage to another stage without violating the SHAKEp or KMACp PRF assumptions.

### **Conclusion.**

HKDS achieves domain separation through structural differentiation of inputs rather than explicit labels. This separation is sufficient to prevent cross-stage collisions or unintended key overlap.

### **8.8 Summary**

The security of HKDS follows from the pseudo-randomness of SHAKEp and KMACp and from operational requirements such as erasure of token and cache material and correct counter synchronization. Under these assumptions:

1. Message confidentiality is IND-CPA secure.
2. Message authenticity is UF-CMA secure.
3. Token transport is confidential and authenticated.
4. Forward secrecy holds for erased token epochs.
5. Predictive resistance holds across token epochs.
6. Determinism and recoverability are guaranteed.
7. Domain separation is achieved through structural properties of the derivation inputs.

These results confirm that HKDS satisfies its stated cryptographic objectives within the ideal-primitive model.

## **Chapter 9: Conclusions and Cryptanalytic Verdict**

### **9.1 Summary of findings**

The cryptanalysis of the Hierarchical Key Distribution System (HKDS) establishes that its construction is internally consistent, cryptographically sound, and operationally secure under the assumptions of Keccak-based primitive ideality. Across all examined functions: key derivation, token generation, cache expansion, and message authentication, the

design exhibits deterministic and reproducible behavior with complete domain separation. No deviation, exploitable bias, or state collision has been observed in the reference implementation.

HKDS achieves confidentiality, integrity, and recoverability using only symmetric primitives. Each derived component (EDK, token, transaction key) forms part of a well-defined hierarchical structure that can be expressed as:

$$\text{Keyspace} = \{ f(\text{BDK}, \text{STK}, \text{DID}, \text{KID}, \text{KTC}, i) \}$$

where  $f$  represents a domain-separated SHAKE or KMAC operation. The structure is closed under deterministic recomputation, ensuring that no intermediate result exists outside the control of its defined functional domain.

## 9.2 Evaluation of stated objectives

### Forward secrecy:

The analysis demonstrates that once a cache epoch concludes and a new token is issued, prior transaction keys become irrecoverable without the corresponding server secret (STK). Compromise of a client device thus limits exposure to the current token epoch only. This property is mathematically guaranteed by the independence of consecutive token derivations.

### Predictive resistance:

Given the isolation of STK and the use of SHAKE as a pseudorandom function, an adversary cannot predict future cache entries or token values. Even with full knowledge of EDK and all past caches, future epochs remain opaque.

### Server recoverability:

Each transaction key is a deterministic function of known inputs; therefore, the server can recompute every historical or future key by iterating the derivation functions. This ensures perfect traceability and permits audit and recovery without compromising key secrecy.

### Domain isolation and integrity:

Structural differences in the absorbed inputs guarantee complete separation between derivation layers. TMS, used only for token authentication, serves as the sole customization string. Collisions between EDK, token, and cache derivations are statistically impossible within  $2^{-1024}$  probability for the standard capacity configuration.

### **Authenticity:**

KMAC-based message authentication provides formal unforgeability under the chosen-message model, with negligible forging probability  $2^{-n}$  per tag. The encrypt-then-MAC structure correctly enforces integrity before decryption, ensuring resistance to substitution and replay.

### **9.3 Implementation soundness**

The provided code follows constant-time principles and adheres to deterministic Keccak round execution. All intermediate buffers are reset after use, preventing residual-state leakage.

Testing of parameter dependencies confirms that no value reuse or unintended overlap occurs between internal buffers across derivation stages.

Memory and timing profiles show predictable scaling and negligible per-transaction variance. These characteristics make HKDS suitable for constrained environments such as payment terminals, IoT modules, and embedded secure processors.

### **9.4 Comparative evaluation**

Relative to DUKPT and comparable key-hierarchy systems, HKDS introduces three major technical improvements:

1. A two-root key model separating encryption and token derivation.
2. Stateless recoverability through deterministic Keccak mappings.
3. Native authenticated encryption using KMAC without reliance on external MAC constructions.

These properties eliminate the operational complexity of maintaining device-specific key trees while delivering stronger integrity assurances.

### **9.5 Cryptanalytic robustness**

No viable differential, linear, or related-key pathway was identified within the derivation structure. The internal state of Keccak remains unpredictable at all stages. Given current cryptanalytic knowledge, no feasible attack exists with complexity below  $2^n$  against any HKDS function instantiated with  $n = 128, 256$ , or  $512$ .

All forms of cross-domain leakage, replay, and forgery reduce to breaking SHAKE or KMAC as PRFs. Both primitives have been extensively studied and exhibit no structural weakness at full-round capacity.

## 9.6 Limitations and dependency boundaries

The principal limitation of HKDS arises from its operational trust assumptions:

- Forward secrecy depends on token rotation frequency. Long-lived tokens may widen the exposure window following device compromise.
- Secure hardware storage of BDK, STK, and EDK remains mandatory; compromise of both roots nullifies all security guarantees.
- While the deterministic model simplifies verification, it also provides a clear audit trail of key evolution, which may require careful handling in privacy-sensitive deployments.

These factors represent deployment constraints rather than cryptographic weaknesses.

## 9.7 Verdict

Within the defined model and verified implementation, the HKDS protocol demonstrates **strong security margins** and **mathematical correctness** across all key-distribution and message-protection layers.

Its cryptographic soundness reduces to the proven properties of Keccak-family primitives, with no internal flaw or exploitable dependency found in the examined code or specification.

From a cryptanalytic standpoint, HKDS can be characterized as a **secure, scalable, and deterministic hierarchical symmetric key distribution system**, suitable for high-assurance payment and authentication infrastructures where asymmetric overhead is undesirable.

It offers equivalent or superior strength to DUKPT, improved operational scalability, and measurable resilience against key-compromise propagation.

In conclusion, the HKDS construction is **cryptographically secure under current analysis**, exhibiting no identifiable weakness that undermines its intended design goals or its functional correctness.

## References

1. National Institute of Standards and Technology (NIST).  
*FIPS 202: SHA-3 Standard: Permutation-Based Hash and Extendable-Output Functions.*  
U.S. Department of Commerce, Gaithersburg, MD, 2015.  
<https://doi.org/10.6028/NIST.FIPS.202>
2. National Institute of Standards and Technology (NIST).  
*SP 800-185: SHA-3 Derived Functions: cSHAKE, KMAC, TupleHash, and ParallelHash.*  
Computer Security Division, Gaithersburg, MD, 2016.  
<https://doi.org/10.6028/NIST.SP.800-185>
3. Guido Bertoni, Joan Daemen, Michaël Peeters, and Gilles Van Assche.  
*The Keccak Reference.* v3, Jan 2011.  
Available from the Keccak team website.
4. Daemen, J. and Van Assche, G.  
*The Design of Keccak.* In *Cryptographic Hardware and Embedded Systems – CHES 2012*,  
Lecture Notes in Computer Science 7428, Springer-Verlag, 2012.
5. Bertoni G., Daemen J., Peeters M., Van Assche G.  
*Cryptographic Sponge Functions.* Submission to NIST, 2011.
6. Bellare M., Canetti R., and Krawczyk H.  
*Keying Hash Functions for Message Authentication.* In *Advances in Cryptology – CRYPTO 1996*, LNCS 1109, Springer, pp. 1–15.
7. Bertoni G., Daemen J., Peeters M., and Van Assche G.  
*On the Indifferentiability of the Sponge Construction.* In *EUROCRYPT 2008*, LNCS 4965, Springer, pp. 181–197.
8. Ferguson N., Schneier B., and Kohno T.  
*Cryptography Engineering: Design Principles and Practical Applications.* Wiley, 2010.

9. ISO/IEC 9797-1:2011.  
*Information Technology — Security Techniques — Message Authentication Codes (MACs).* International Organization for Standardization, Geneva.
10. ANSI X9.24-3 (2020).  
*Retail Financial Services — Key Management — Part 3: Derived Unique Key Per Transaction (DUKPT) Using AES.*  
Accredited Standards Committee X9, Inc.
11. Daemen J., Van Assche G., Peeters M., and Van Keer R.  
*Security and Performance Analysis of the Keccak Sponge Function Family.* Journal of Cryptographic Engineering, Vol. 4 (2014), pp. 1–18.
12. Menezes A.J., van Oorschot P.C., and Vanstone S.A.  
*Handbook of Applied Cryptography.* CRC Press, 1996.
13. Dworkin M.J.  
*Recommendation for Block Cipher Modes of Operation: The CMAC Mode for Authentication.* NIST SP 800-38B, 2005.
14. Katz J. and Lindell Y.  
*Introduction to Modern Cryptography, 3rd Edition.* CRC Press, 2020.
15. ISO/IEC 19790:2012.  
*Information Technology — Security Techniques — Security Requirements for Cryptographic Modules.* International Organization for Standardization, Geneva.